

# Computer Architectures

## Chapter 4

Tien-Fu Chen

National Chung Cheng Univ.

© by Tien-Fu Chen@CCU

chap4-0

### Advance Pipelining

#### ☐ Static Scheduling

Have compiler to minimize the effect of structural, data, and control dependence

- **advantages:** simple hardware

- **Examples:**

Loop unrolling

Software Pipelining

Trace Scheduling

#### ☐ Dynamic Scheduling

Have hardware to rearrange instruction execution to reduce the stalls

- **advantages:**
  - handle dependence unknown at compile time
  - simplify compiler design
  - code compatible

- **Examples:**

Scoreboarding

Tomasulo's Algorithm

Dynamic Branch Prediction

© by Tien-Fu Chen@CCU

chap4-1

# Instruction Level Parallelism

## ❑ Taking advantages of ILP

- **Superscalar** - multiple issue per cycle
- **Superpipelining** - deeper pipelines
- **VLIW** - very long instruction word

## ❑ Increase ILP by compiler

- **Loop unrolling** -  
Increase instructions between loop branch by replicating loop body multiple times
- **Software Pipelining**  
Reorganize loop code such that each iteration contains code chosen from different iterations
- **Trace Scheduling**  
Increase parallelism by selecting more code candidates

## ❑ Usually hardware techniques require compiler support

- Superscalar needs instruction scheduling
- VLIW needs trace scheduling

# Static Scheduling - Loop Unrolling

## ❑ Basic Idea

```
loop:  LD    F0, 0(R1)
      ADD   F4, F0, F2
      SUB   R1, R1, #8
      BNE   R1, loop
      SD    0(R1), F4
```

```
loop:  LD    F0, 0(R1)
      ADD   F4, F0, F2
      SD    0(R1), F4
```

## ❑ Features


- reduce loop overhead
  - branch
  - counter
- increase code size
- increase register requirement

```
LD    F6, 8(R1)
ADD   F8, F6, F2
SD    8(R1), F8
```

```
LD    F10, 16(R1)
ADD   F12, F10, F2
SD    16(R1), F12
```

```
LD    F14, 24(R1)
ADD   F16, F14, F2
SD    24(R1), F16
```

```
SUB   R1, R1, #32
BNE   R1, loop
```



# **Identify Dependencies by Compiler**

## **□ Goal:**

compiler concerned about dependencies in program, whether or not a HW hazard depends on a given pipeline

- Data dependency (True)
- Name dependency (anti-/output dependency)
- Control dependency

## **□ (True) Data dependencies**

(RAW if a hazard for HW)

- Instruction i produces a result used by instruction j, or
- Instruction j is data dependent on instruction k, and instruction k is data dependent on instruction i.
- Easy to determine for registers (fixed names)
- Hard for memory:
  - Does  $100(R4) = 20(R6)$ ?
  - From different loop iterations, does  $20(R6) = 20(R6)$ ?

# **Name dependency**

## **□ Name dependence:**

two instructions use same name but don't exchange data

- Antidependence (WAR if a hazard for HW)

Instruction j writes a register or memory location that instruction i reads from and instruction i is executed first
- Output dependence (WAW if a hazard for HW)

Instruction i and instruction j write the same register or memory location; ordering between instructions must be preserved.
- no value being transmitted between instructions

## **□ Register Renaming**

- Statically by compiler
- Dynamic by hardware

# **Control dependency**

## **□ control dependence**

Example

```
if p1 {S1;};  
if p2 {S2;}
```

- S1 is control dependent on p1 and
- S2 is control dependent on p2 but not on p1.

## **□ Two constraints on control dependencies:**

- An instruction that is control dependent on a branch cannot be moved before the branch so that its execution is no longer controlled by the branch.
- An instruction that is not control dependent on a branch cannot be moved to after the branch so that its execution is controlled by the branch.

## **□ Control dependencies relaxed to get parallelism; get same effect if**

- preserve exception behavior
- preserve data flow

# **Loop-carried dependence**

## **□ Where are data dependencies?**

```
for (i=1; i<=100; i=i+1)  
{  
    A[i+1] = A[i] + C[i];    /* S1 */  
    B[i+1] = B[i] + A[i+1]; /* S2 */  
}
```

1. S2 uses the value, A[i+1], computed by S1 in the same iteration.
2. S1 uses a value computed by S1 in an earlier iteration, since iteration i computes A[i+1] which is read in iteration i+1. The same is true of S2 for B[i] and B[i+1].

## **□ loop-carried dependence**

dependency exists between different iterations of the loop

- Loop with loop-carried dependence cannot be executed in parallel

## **□ GCD test :**

```
write A[a * i+b]  
read A[c * i+d]
```

if loop-carried depend exists then

$$(d - b) \bmod \text{GCD}(a,c) = 0$$

## Now Safe to Unroll Loop? (p. 240)

OLD:                   for (i=1; i<=100; i=i+1) {  
                          A[i+1] = A[i] + B[i];   /\* S1 \*/  
                          B[i+1] = C[i] + D[i];} /\* S2 \*/

A[1] = A[1] + B[1];

NEW:                   **for**       (i=1; i<=99; i=i+1) {  
                          B[i+1] = C[i] + D[i];  
                          A[i+1] = + A[i+1] + B[i+1];  
                          }  
  
                          B[101] = C[100] + D[100];

© by Tien-Fu Chen@CCU

chap4-8

## Software Pipelining

### □ Software pipelining


Reorganize loops such that each iteration contains instruction sequences from different iterations in original code

- need start-up & finish blocks

1. read	1. read
2.	2.       read
3. add	3. add       read
4. write	4. write   add       read
	5.       write   add
	6.       write   add
	7.       write

<b>for</b> i = 1 , n	<b>for</b> i = 4 , n
<i>read<sub>i</sub></i>	<i>write<sub>i-3</sub></i>
<i>add<sub>i</sub></i>	<i>add<sub>i-1</sub></i>
<i>write<sub>i</sub></i>	<i>read<sub>i</sub></i>
<b>end</b>	<b>end</b>



## Compared with loop unrolling

### Before: Unrolled 3 times


```

1  L.D    F0,0(R1)
2  ADD.D  F4,F0,F2
3  S.D    0(R1),F4
4  L.D    F6,-8(R1)
5  ADD.D  F8,F6,F2
6  S.D    -8(R1),F8
7  L.D    F10,-16(R1)
8  ADD.D  F12,F10,F2
9  S.D    -16(R1),F12
10 DSUBUI R1,R1,#24
11 BNEZ   R1,LOOP
    
```

### After: Software Pipelined

```

1  S.D    0(R1),F4 ; Stores M[i]
2  ADD.D  F4,F0,F2 ; Adds to M[i-1]
3  L.D    F0,-16(R1); Loads M[i-2]
4  DSUBUI R1,R1,#8
5  BNEZ   R1,LOOP
    
```



### • Symbolic Loop Unrolling

- Maximize result-use distance
- Less code space than unrolling
- Fill & drain pipe only once per loop vs. once per each unrolled iteration in loop unrolling

## Trace Scheduling

### ❑ Trace



a most likely sequence of instructions

### ❑ Two phases

- **trace selection**
  - identify frequent codes
  - may use loop unrolling to generate long trace
  - should consider data dependence constraints and branch points
- **trace compaction**
  - squeeze trace into a small of wide instructions
  - move instruction as early as possible
  - compact instructions as few as possible
  - add compensation code

### ❑ good or bad?

- scheduling across basic blocks
- code explosion



## **Summary of Static Scheduling**

### **❑ Loop unrolling**

- + multiple loop body for scheduling
- + reduce branch frequency
- expand code size
- must handle "residual" iterations
- increase register pressure

### **❑ Software Pipelining**

- + no dependences in loop body
- does not reduce branch hazards
- need start-up and finish blocks
- increase register pressure

### **❑ Trace Scheduling**

- + works for loops
- + increase most likely operations for VLIW
- more complex than loop unrolling
- code expansion

© by Tien-Fu Chen@CCU

chap4-12

## **Advance Pipelining again**

### **❑ Static Scheduling**

Have compiler to minimize the effect of structural, data, and control dependence

### **❑ Dynamic Scheduling**

HW rearrange the instruction execution to reduce stalls

### **❑ Key idea**

Allow instructions behind stall to proceed

**DIVD F0,F2,F4**

**ADDD F10,F0,F8**

**SUBD F8,F8,F14**

- Enables out-of-order execution => out-of-order completion
- ID stage checked both for structural & data dependencies
- Splitting ID into two stages
  - Issue - decode instructions, check for structural hazards
  - Read operands - wait until no data hazards, then read operands

© by Tien-Fu Chen@CCU

chap4-13

## **Advantages of Dynamic Scheduling**

- ☐ **Handles cases when dependences unknown at compile time (e.g., because they may involve a memory reference)**
- ☐ **It simplifies the compiler**
- ☐ **Allows code that compiled for one pipeline to run efficiently on a different pipeline**
- ☐ **Hardware speculation, a technique with significant performance advantages, that builds on dynamic scheduling**

## **Out-of-order execution**

### ☐ **Split pipelining stages:**

- Original DLX pipelining
- New DLX pipelining

IF -

ID - decode, check for hazards, fetch operands

EX - execute instruction

WB -

**Issue -** decode, check structural hazards

read operands - wait until no data hazards

execute -

write back-

### ☐ **Scoreboarding**

- multiple function units
- each instruction goes through the scoreboard
- centralized control scheme
  - control all instruction issue
  - detect all hazards
  - maintain hazard resolutions
- implemented in CD 6600



# **Scoreboarding**

## **❑ Dealing with hazards**

- **Issue:** check structural hazards  
check for WAW hazards  
stall issue until hazard cleared
- **Read operand:** check for RAW hazards  
wait until data ready
- **Execution:** execute operations  
notify scoreboard when completion
- **Write back:** check for WAR  
stall write until clear

## **❑ Maintain three data structures**

- **Instruction status -**  
indicate the four steps of instruction in
- **Function unit status -**  
indicate states of each function units
- **Register result status -**  
indicate which function unit will write a register

# **Three Parts of the Scoreboard**

1. **Instruction status**—which of 4 steps the instruction is in
2. **Functional unit status**—Indicates the state of the functional unit (FU). 9 fields for each functional unit
  - Busy**—Indicates whether the unit is busy or not
  - Op**—Operation to perform in the unit (e.g., + or −)
  - Fi**—Destination register
  - Fj, Fk**—Source-register numbers
  - Qj, Qk**—Functional units producing source registers Fj, Fk
  - Rj, Rk**—Flags indicating when Fj, Fk are ready
3. **Register result status**—Indicates which functional unit will write each register, if one exists. Blank when no pending instructions will write that register

## Detailed control and data maintaining

Instruction status	Wait until	Bookkeeping
Issue	Not busy (FU) and not result(D)	$\text{Busy}(\text{FU}) \leftarrow \text{yes}; \text{Op}(\text{FU}) \leftarrow \text{op}; \text{Fi}(\text{FU}) \leftarrow 'D';$ $\text{Fj}(\text{FU}) \leftarrow 'S1'; \text{Fk}(\text{FU}) \leftarrow 'S2';$ $\text{Qj} \leftarrow \text{Result}('S1'); \text{Qk} \leftarrow \text{Result}('S2'); \text{Rj} \leftarrow \text{not } \text{Qj}; \text{Rk} \leftarrow \text{not } \text{Qk}; \text{Result}('D') \leftarrow \text{FU};$
Read operands	Rj and Rk	$\text{Rj} \leftarrow \text{No}; \text{Rk} \leftarrow \text{No}$
Execution complete	Functional unit done	
Write result	$\forall f((\text{Fj}(f) \neq \text{Fi}(\text{FU}) \text{ or } \text{Rj}(f) = \text{No}) \& (\text{Fk}(f) \neq \text{Fi}(\text{FU}) \text{ or } \text{Rk}(f) = \text{No}))$	$\forall \ell(\text{if } \text{Qj}(\ell) = \text{FU} \text{ then } \text{Rj}(\ell) \leftarrow \text{Yes});$ $\forall \ell(\text{if } \text{Qk}(\ell) = \text{FU} \text{ then } \text{Rk}(\ell) \leftarrow \text{Yes});$ $\text{Result}(\text{Fi}(\text{FU})) \leftarrow 0; \text{Busy}(\text{FU}) \leftarrow \text{No}$

### □ Try to identify how scoreboard deals with

- RAW
- WAR
  - Queue both the operation and copies of its operands
  - Read registers only during Read Operands stage
- WAW

© by Tien-Fu Chen@CCU

chap4-18

		Instruction status			
Instruction		Issue	Read operands	Execution complete	Write result
LD	F6, 34(R2)	√	√	√	√
LD	F2, 45(R3)	√	√	√	
MULTD	F0, F2, F4	√			
SUBD	F8, F6, F2	√			
DIVD	F10, F0, F6	√			
ADDD	F6, F8, F2				

Functional unit status									
Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
Integer	Yes	Load	F2	R3				No	
Mult1	Yes	Mult	F0	F2	F4	Integer		No	Yes
Mult2	No								
Add	Yes	Sub	F8	F6	F2		Integer	Yes	No
Divide	Yes	Div	F10	F0	F6	Mult1		No	Yes

Register result status									
	F0	F2	F4	F6	F8	F10	F12	...	F30
FU	Mult1	Integer			Sub	Divide			

## **Scoreboarding (cont)**

### **❑ Benefits**

- Improvement of 1.7 for Fortran programs
- 2.5 for hand-coded assembly program
- How about modern compiler??

### **❑ Cost**

- only as much as one function unit
- main cost in buses ( 4 times)

### **❑ Limitations**

- Results are through register file, never data forwarded.
- Data dependency must wait for access to register file
- Stalls when WAW hazards occur
- Limited to instructions in basic block (small window)

## **Tomasulo's Algorithm**

### **❑ Key features:**

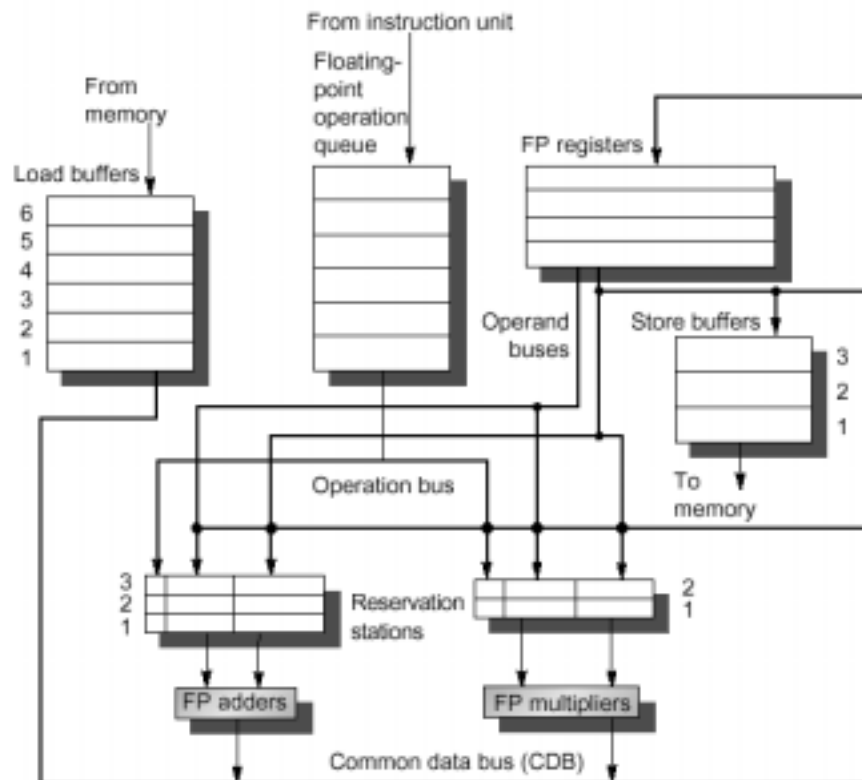
- Reservation stations at each functional unit
- A Common Data Bus (CDB) to broadcast all results
- Employ register renaming
- Use "tag" to handle hazard control
- load/store buffers

### **❑ Only three steps (not include IF, MEM)**

- Issue -get instruction  
    check available function unit  
    or check available load buffer  
    stalls on structural hazards
- Execute - execute when operands available  
    if not, check CDB for operand
- Write back - if CDB available, write result  
    if not, stalls on CDB

### **❑ Data structures are attached to reservation stations, reg file, load buffer**

## Data structure of Tomasulo's Algorithm



© by Tien-Fu Chen@

chap4-22

## Reservation Station Components

**Op**—Operation to perform in the unit (e.g., + or −)

**V<sub>j</sub>, V<sub>k</sub>**—Value of Source operands

- Store buffers has V field, result to be stored

**Q<sub>j</sub>, Q<sub>k</sub>**—Reservation stations producing source registers (value to be written)

- Note: No ready flags as in Scoreboard; Q<sub>j</sub>, Q<sub>k</sub>=0 => ready
- Store buffers only have Q<sub>i</sub> for RS producing result

**Busy**—Indicates reservation station or FU is busy

**Register result status**—Indicates which functional unit will write each register, if one exists. Blank when no pending instructions that will write that register.

# Bookkeeping rules

Instruction status	Wait until	Action or bookkeeping
Issue	Station or buffer empty	<pre> if (Register['S1'].Qi ≠ 0)     {RS[r].Qj ← Register['S1'].Qi} else {RS[r].Vj ← S1; RS[r].Qj ← 0}; if (Register[S2].Qi ≠ 0)     {RS[r].Qk ← Register[S2].Qi}; else {RS[r].Vk ← S2; RS[r].Qk ← 0} RS[r].Busy ← yes; Register['D'].Qi = r; </pre>
Execute	(RS[r].Qj=0) and (RS[r].Qk=0)	None—operands are in Vj and Vk
Write result	Execution completed at r and CDB available	<pre> Vx{if (Register[x].Qi=r) {Fx ← result;     Register[x].Qi ← 0}}; Vx{if (RS[x].Qj=r) {RS[x].Vj ← result;     RS[x].Qj ← 0}}; Vx{if (RS[x].Qk=r) {RS[x].Vk ← result;     RS[x].Qk ← 0}}; Vx{if (Store[x].Qi=r) {Store[x].V ← result;     Store[x].Qi ← 0}}; RS[r].Busy ← No </pre>

## □ How to eliminating WAW and WAR?

© by Tien-Fu Chen@CCU

chap4-24

Instruction status			
Instruction		Issue	Execute      Write result
LD      F6, 34(R2)		✓	✓      ✓
LD      F2, 45(R3)		✓	✓
MULTD   F0, F2, F4		✓	
SUBD    F8, F6, F2		✓	
DIVD    F10, F0, F6		✓	
ADDD    F6, F8, F2		✓	

Reservation stations						
Name	Busy	Op	Vj	Vk	Qj	Qk
Add1	Yes	SUB	Mem[34+Regs[R2]]			Load2
Add2	Yes	ADD			Add1	Load2
Add3	No					
Mult1	Yes	MULT		Regs[F4]	Load2	
Mult2	Yes	DIV		Mem[34+Regs[R2]]	Mult1	

Register status									
Field	F0	F2	F4	F6	F8	F10	F12	...	F30
Qi	Mult1	Load2		Add2	Add1	Mult2			

Instruction		Instruction status		
		Issue	Execute	Write result
LD	F6, 34 (R2)	✓	✓	✓
LD	F2, 45 (R3)	✓	✓	✓
MULTD	F0, F2, F4	✓	✓	
SUBD	F8, F6, F2	✓	✓	✓
DIVD	F10, F0, F6	✓		
ADDD	F6, F8, F2	✓	✓	✓

Reservation stations						
Name	Busy	Op	Vj	Vk	Qj	Qk
Add1	No					
Add2	No					
Add3	No					
Mult1	Yes	MULT	Mem[45+Regs[R3]]	Regs[F4]		
Mult2	Yes	DIV		Mem[34+Regs[R2]]	Mult1	

Register status									
Field	F0	F2	F4	F6	F8	F10	F12	...	F30
Qi	Mult1					Mult2			

© by

.26

## With load/store buffers

Instruction status				
Instruction	From iteration	Issue	Execute	Write result
LD F0, 0(R1)	1	✓	✓	
MULTD F4, F0, F2	1	✓		
SD 0(R1), F4	1	✓		
LD F0, 0(R1)	2	✓	✓	
MULTD F4, F0, F2	2	✓		
SD 0(R1), F4	2	✓		

Reservation stations						
Name	Busy	Op	Vj	Vk	Qj	Qk
Add1	No					
Add2	No					
Add3	No					
Mult1	Yes	MULT		Regs[F2]	Load1	
Mult2	Yes	MULT		Regs[F2]	Load2	

Register status									
Field	F0	F2	F4	F6	F8	F10	F12	...	F30
Qi	Load2		Mult2						

Load buffers			
Field	Load 1	Load 2	Load 3
Address	Regs[R1]	Regs[R1]-8	
Busy	Yes	Yes	No

Store buffers			
Field	Store 1	Store 2	Store 3
Qi	Mult1	Mult2	
Busy	Yes	Yes	No
Address	Regs[R1]	Regs[R1]-8	

© by Tien-Fu Chen@CC

.27

# **Tomasulo's Algorithm (cont)**

## ❑ Differences from scoreboarding

- Control & buffers distributed with Function Units,
- HW renaming of registers to avoid WAW and WAR hazards
- CDB to broadcast results rather than waiting on registers
- load/store buffers are treated as function unit

## ❑ Reservation Stations

- Handle distributed hazard detection and instruction control
- Use 4-bit tag field to specify which RV station or load buffer will produce result

## ❑ Register renaming

- Tag assigned on instruction issue
- Tag discarded after write back to register
- WAW and WAR are eliminated
- Even better if a branch-taken strategy on a loop  
    ➡ multiple execution simultaneously  
    loop is unrolled dynamically by hardware

# **Comparisons**

## **Scoreboarding**



- Global data structure  
  centralized control
- Use register designator  
  + simple, low cost
- Structural stalls on FU
- Solve :  
  RAW by register  
  WAR in write  
  WAW stalls on issue
- Limited: ➡  
  register serialized  
  input      output
- Stalls issue on block  
     $S1 \leftarrow S2 * S3$   
     $S4 \leftarrow S5 * S6$

## **Tomasulo's Alg**

- Distributed control
- Use tagged register  
  + register renaming  
  - tag allocation and  
  deallocation  
  - associative compare
- Structural stalls on RS
- Solve:  
  RAW by CDB  
  WAR copy opnd to RS  
  WAW by renaming
- Limited: CDB  
  + broadcast result  
  - one result per cycle

# Dynamic Branch Prediction

## ❑ For a loop




## ❑ Branch History Table

- Lower bits of PC address index table of 1-bit values
- Tells whether or not branch taken last time
- Problems:
  - the end of a loop and 1st time of next loop

## ❑ 2-bit scheme

where change prediction only if get misprediction twice.



© by Tien-Fu Chen@CCU

chap4-30

# Performance of Prediction Buffer

## ❑ Mispredict because either:

- Wrong guess for that branch
- Got history of wrong branch when index table

## ❑ Prediction Accuracy

- 4096 entry table programs vary from 1% misprediction (nasa7, tomcatv) to 18% (eqntott), with spice at 9% and gcc at 12%
- 4096 about as good as infinite table, but 4096 is a lot of HW

## ❑ Performance of prediction buffer:

Given that 90% hit in buffer and 90% of correct guess

$$\begin{aligned} \text{Accuracy} &= (\% \text{hit in buffer} * \% \text{correct}) + \\ &\quad [(1 - \% \text{hit in buffer}) * \text{luck guess}] \\ &= (90\% \times 90\%) + (50\% * 10\%) \\ &= 86\% \end{aligned}$$

© by Tien-Fu Chen@CCU

chap4-31



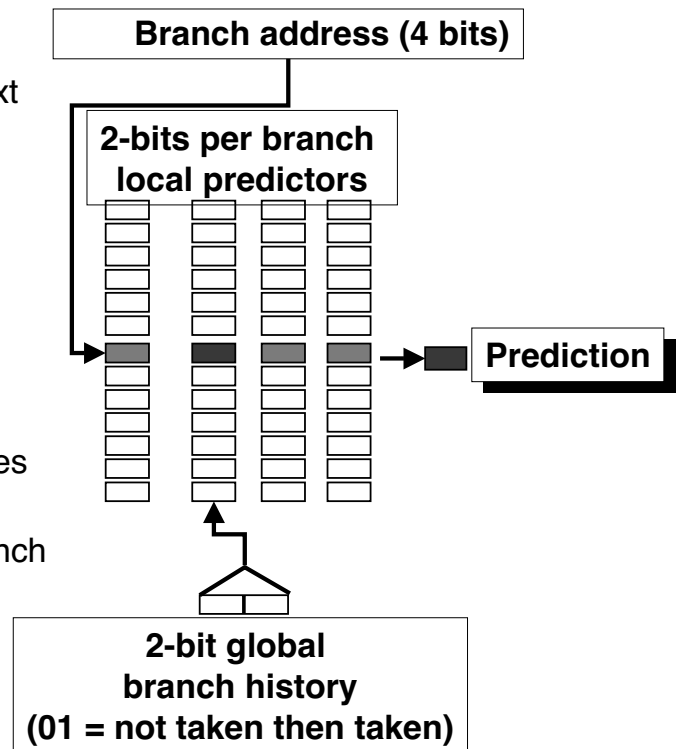
# Correlating Branch Predictor

## ❑ Idea:

taken/not taken of recently executed branches is related to behavior of next branch  
(as well as the history of that branch behavior)

## ❑ a (m,n) predictor

- use the behavior of the last m branches
- choose from  $2^m$  branch predictors
- each is n-bit predictor for a single branch
- cost:  
 $2^m \times n \times \# \text{ of prediction entries selected by addr}$



# Branch Target Buffer

## ❑ Guess both branch cond and branch target

IF: check BTB  
ID: check cond  
EX: refill BTB

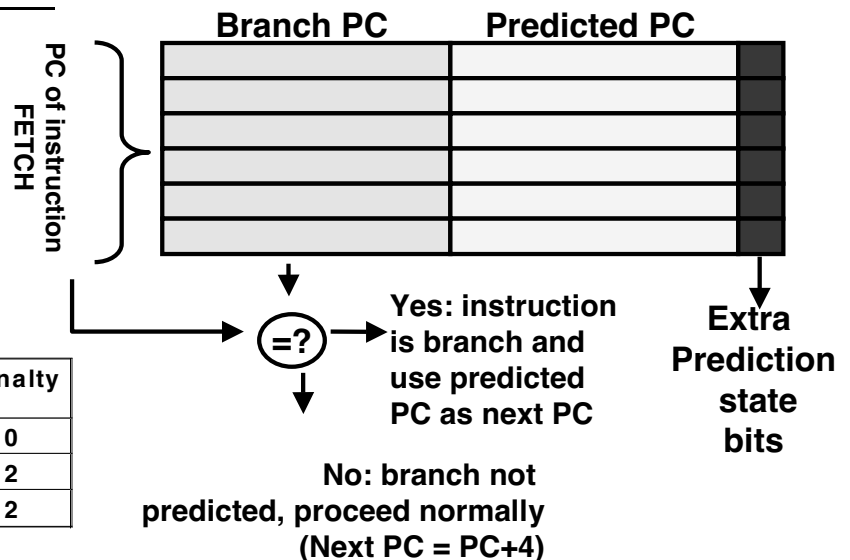
## ❑ Penalty

hit in buffer	prediction	actual branch	Penalty
Yes	Taken	Taken	0
Yes	Taken	N Taken	2
No		Taken	2

## ❑ Total penalty

(60% taken branch)

$$\begin{aligned}
 \text{Penalty} &= \% \text{ hit in buffer} * \% \text{ incorrect} * 2 + \\
 &\quad (1 - \% \text{ hit in buffer}) * \% \text{ taken} * 2 \\
 &= (90\% * 10\% * 2) + (10\% * 60\% * 2) \\
 &= 0.30 \text{ cycles}
 \end{aligned}$$



## **Variations of BTB**

### **❑ Separating target buffer from a prediction buffer - PowerPC 620**

- target address prediction
- branch direction

### **❑ Storing target instructions**

- allow longer time to access BTB, better for a larger BTB
- allowed to perform branch folding

### **❑ Predicting indirect jumps**

destination address varies at run-time

e.g., indirect procedure calls, procedure returns,

- return address stack

### **❑ Reducing misprediction penalty**

- fetching from both predicted and unpredicted direction ->requiring dual-ported instruction memory
- caching instructions from multiple paths

## **Multiple-Issue Processors**

### **❑ Superscalar**

- Issue varying # of instructions per clock
  - be either statically scheduled by compiler or dynamically scheduled using hardware
- e.g. IBM PowerPC, Sun SuperSparc, DEC Alpha, HP 7100

### **❑ Very Long Instruction Words (VLIW)**

- fixed number of instructions scheduled by the compiler
- inherently statically scheduled by compiler
- Joint HP/Intel agreement in 1998??

### **❑ Hardware support for more ILP**

- Conditional/Predicated Instructions
- Speculative instructions with renaming
- Speculative execution

# **Superscalar**

## ❑ **Superscalar machines**

- Issue multiple **independent** instructions per clock cycle
- if dependence exist, only first instruction issued
- hardware not dynamically concern about issue
- need  $n$  function units for degree of  $n$
- would help if compiler properly schedule codes

## ❑ **Limitation**

- contention on FP registers  
why: FP load/store vs FP operations  
solution: dual-ported register file?
- Delayed load  
why: original one delay slot, now 3 slots
- Delayed branch

## ❑ **Advantages:**

- code compatibility
- could reach  $CPI < 1$

© by Tien-Fu Chen@CCU

chap4-36

# **Dynamic Scheduling in Superscalar**

## ❑ **Dependencies stop instruction issue**

- Code compiler for scalar version will run poorly on superscalar
- May want code to vary depending on how superscalar

## ❑ **Simple approach:**

- separate Tomasulo Control for separate reservation stations for Integer FU/Reg and for FP FU/Reg
- Issue 2X Clock Rate, so that issue remains in order
- Only FP loads might cause dependency between integer and FP issue:
  - Replace load reservation station with a load queue; operands must be read in the order they are fetched
  - Load checks addresses in Store Queue to avoid RAW violation
  - Store checks addresses in Load Queue to avoid WAR,WAW

## **VLIW (Very Long Instruction Word)**

### **❑ tradeoff instruction space for simple decoding**

- long instruction word has room for many operations
- all the operations the compiler puts in the long instruction word can execute in parallel

E.g., 2 int ops, 2 FP ops, 2 mems, 1 branch

16 to 24 bits per field  $\Rightarrow 7 \times 16 = 112$  bits to  $7 \times 24 = 168$  bits wide

### **❑ Features :**

- Central controller issues a single long instruction
- Each instruction initiates many independent operations
- Each operation executed in fixed cycles
- Operations can be pipelined

### **❑ Compiler Aid**

- loop unrolling
- scheduling code across basic blocks

**trace scheduling**

## **Limitations on multiple-issue**

### **❑ Limited ILP**

- 1 branch in 5 instructions  $\Rightarrow$  how to keep a 5-way VLIW busy?
- Latencies of units  $\Rightarrow$  many operations must be scheduled
- Need Pipeline Depth  $\times$  No. Functional Units of independent operations to keep machines busy

### **❑ Limits on VLIW**

- Limited amount of parallelism available in instruction sequences
- require a large-number of memory and register bandwidth for different functional units at the same time
- code size explosion  
loop unrolling + no op code
- VLIW lock step  $\Rightarrow$  1 hazard & all instructions stall
- binary compatibility is practical weakness

## Comparisons

### Superscalar

### VLIW

#### adv. :

- better code density
- code compatible

#### difference :

- issue decision at: Run Time

#### disadv :

- require more instruction fetching & decoding
- more delayed slot to be filled
- different functional units
- multi-ported reg. file
- limited parallelism

#### adv. :

- fixed instruction format
- more parallelism provided

- by trace scheduling at : Compile Time

#### disadv :

- static scheduling - compiler
- no dynamic decision
- code explosion
- multi-ported register file
- limited parallelism
- difficult to use data cache because of sharing

## Hardware Support for more ILP

### ❑ Conditional (predicate) instructions

- Eliminate branches by turning branches into conditionally executed instructions  
if (x) then A=B op C else NOP
- If false, then neither store result or cause exception
- Expanded ISA of Alpha, MIPS, PowerPC, SPARC have conditional move; PA-RISC can nullify any following instructions

### ❑ Limitation of Conditional Branches

- Still takes a clock even if annulled
- Stall if condition evaluated late
- Limited use when control flow involves more than a simple sequence; Complex conditions reduce effectiveness, condition becomes known late in pipeline
- May increase CPI for conditional instructions or cause a slower overall clock rate

## **Speculation**

### **❑ Speculation:**

Allow an instruction to issue that is dependent on branch predicted to be taken without any consequences (including exceptions) if branch is not actually taken (without undo)

Methods support ambitious speculation

### **❑ HW/SW cooperation**

- HW and OS handle exceptions and return an undefined value for any exception
- may not be acceptable

### **❑ Poison bits**

- a poison bit added to each register and instruction
- a fault occurs on using a value from poison reg

### **❑ Speculation instruction with renaming**

- boosting instruction by flagging as speculation and providing renaming and buffering

© by Tien-Fu Chen@CCU

chap4-42

## **Hardware-based Speculation**

### **❑ Basic idea**

- Allow instruction executed out of order but force them to commit in order; prevent irrevocable action until commits
- separate speculative bypassing of results from real bypassing of results

### **❑ Introducing Reorder Buffer**

- buffer results of uncommitted instructions
- pass results among speculated instructions
- 3 fields: instruction, type, destination, and value
- key differences from Tomasulo's Algorithm:  
Tomasulo: read result from register once result is written  
reorder buffer: supply operands between completion and commit time
- replace load/store buffer
- Use reorder buffer number instead of reservation station buffer


© by Tien-Fu Chen@CCU

chap4-43

## ***Speculative Tomasulo's Algorithm***

## ❑ Four steps of Tomasulo's algorithm

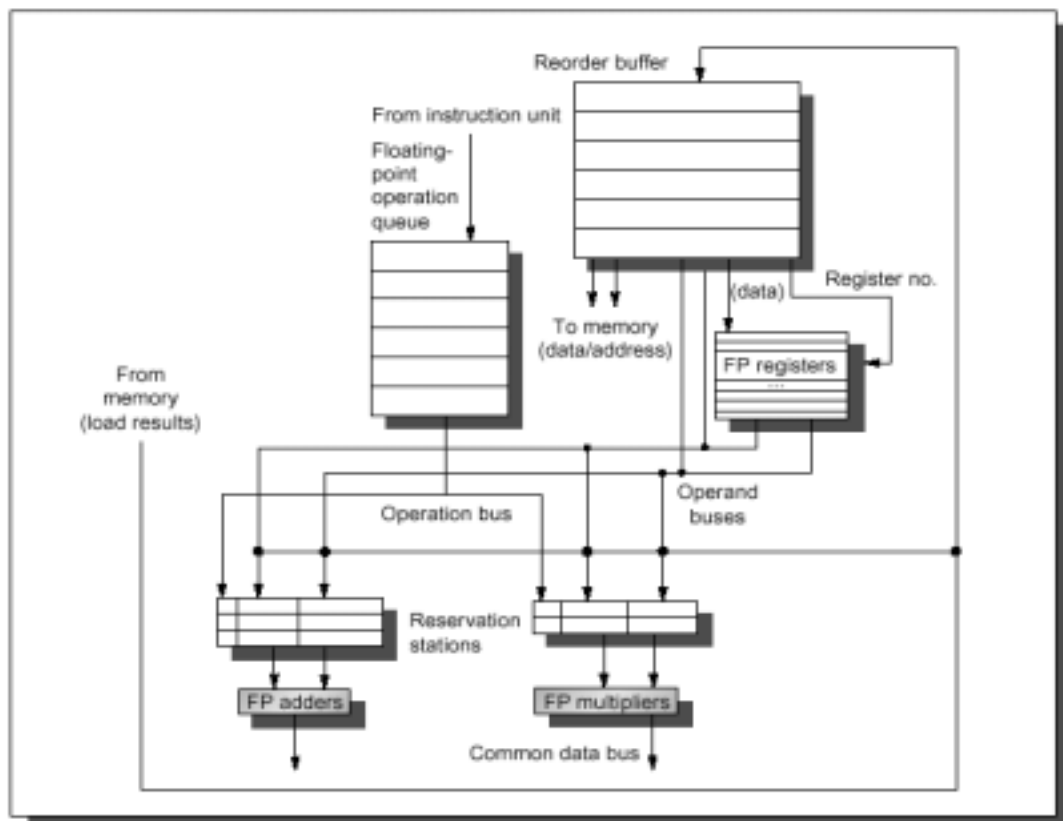
- **Issue**  
If reservation station or reorder buffer slot free, issue instr & send operands & reorder buffer no. for destination.
- **Execution**  
When both operands ready then execute; if not ready, watch CDB for result; when both in reservation station, execute
- **Write result**  
Write on Common Data Bus to all awaiting FUs & reorder buffer; mark reservation station available.
- **Commit**  
When instr. at head of reorder buffer & result present, update register with result (or store to memory) and remove instr from reorder buffer.



© by Tien-Fu Chen@CCU

chap4-44

## Speculative DLX using Tomasulo's Algorithm



© by Tien-Fu C

ap4-45

## **Register renaming, virtual registers versus Reorder Buffers**

- ❑ **Alternative to Reorder Buffer is a larger virtual set of registers and register renaming**
- ❑ **Virtual registers hold both architecturally visible registers + temporary values**
  - replace functions of reorder buffer and reservation station
- ❑ **Renaming process maps names of architectural registers to registers in virtual register set**
  - Changing subset of virtual registers contains architecturally visible registers
- ❑ **Simplifies instruction commit: mark register as no longer speculative, free register with old value**
- ❑ **Adds 40-80 extra registers: Alpha, Pentium,...**
  - Size limits no. instructions in execution (used until commit)

## **Limitation of ILP**

- ❑ **Hardware model of perfect processors**
  - Register renaming  
infinite virtual registers and all WAW & WAR hazards are avoided
  - Branch prediction  
perfect; no mispredictions
  - Jump prediction: jumps perfectly predicted => machine with perfect speculation & an unbounded buffer of instructions available
  - Memory-address alias analysis  
addresses are known & a store can be moved before a load provided addresses not equal
  - 1 cycle latency for all instructions
- ❑ **measurement : IPC**
  - instruction issues per cycle
- ❑ **Observation**
  - FP programs are more sensitive to limit window size



# Workstation Microprocessors 3/2001

Processor	Alpha 21264B	AMD Athlon	HP PA-8600	IBM Power3-II	Intel Pentium III	Intel Pentium 4	MIPS R12000	Sun Ultra-II	Sun Ultra-III
Clock Rate	833MHz	1.2GHz	552MHz	450MHz	1.0GHz	1.5GHz	400MHz	480MHz	900MHz
Cache (I/D/L2)	64K/64K	64K/64K/256K	512K/1M	32K/64K	16K/16K/256K	12K/8K/256K	32K/32K	16K/16K	32K/64K
Issue Rate	4 issue	3 x86 instr	4 issue	4 issue	3 x86 instr	3 x ROPs	4 issue	4 issue	4 issue
Pipeline Stages	7/9 stages	9/11 stages	7/9 stages	7/8 stages	12/14 stages	22/24 stages	6 stages	6/9 stages	14/15 stages
Out of Order	80 instr	72ROPs	56 instr	32 instr	40 ROPs	126 ROPs	48 instr	None	None
Rename regs	48/41	36/36	56 total	16 int/24 fp	40 total	128 total	32/32	None	None
BHT Entries	4K x 9-bit	4K x 2-bit	2K x 2-bit	2K x 2-bit	>= 512	4K x 2-bit	2K x 2-bit	512 x 2-bit	16K x 2-bit
TLB Entries	128/128	280/288	120 unified	128/128	32I / 64D	128I/65D	64 unified	64I/64D	128I/512D
Memory B/W	2.66GB/s	2.1GB/s	1.54GB/s	1.6GB/s	1.06GB/s	3.2GB/s	539 MB/s	1.9GB/s	4.8GB/s
Package	CPGA-588	PGA-462	LGA-544	SCC-1088	PGA-370	PGA-423	CPGA-527	CLGA-787	1368 FC-LGA
IC Process	0.18µ 6M	0.18µ 6M	0.25µ 2M	0.22µ 6m	0.18µ 6M	0.18µ 6M	0.25µ 4M	0.29µ 6M	0.18µ 7M
Die Size	115mm <sup>2</sup>	117mm <sup>2</sup>	477mm <sup>2</sup>	163mm <sup>2</sup>	106mm <sup>2</sup>	217mm <sup>2</sup>	204mm <sup>2</sup>	126 mm <sup>2</sup>	210mm <sup>2</sup>
Transistors	15.4 million	37 million	130 million	23 million	24 million	42 million	7.2 million	3.8 million	29 million
Est mfg cost*	\$160	\$62	\$330	\$110	\$39	\$110	\$125	\$70	\$145
Power(Max)	75W*	76W	60W*	36W*	30W	55W(TDP)	25W*	20W*	65W
Availability	1Q01	4Q00	3Q00	4Q00	2Q00	4Q00	2Q00	3Q0	4Q00

**Max issue: 4 instructions (many CPUs)**

**Max rename registers: 128 (Pentium 4)**

**Max BHT: 4K x 9 (Alpha 21264B), 16Kx2 (Ultra III)**

**Max Window Size (OOO): 126 intructions (Pent. 4)**

**Max Pipeline: 22/24 stages (Pentium 4)**

*Source: Microprocessor Report, www.MPRonline.com*

Processor	Alpha 21264B	AMD Athlon	HP PA-8600	IBM Power 3-II	Intel PIII	Intel P4	MIPS R12000	Sun Ultra-II	Sun Ultra-III
System or Motherboard	Alpha E540 Model 6	AMD GA-72M	HP9000 j6000	RS/6000 44P-170	Dell Prec. 420	Intel D850GB	SGI 2200	Sun Enterprs 450	Sun Blade 1000
Clock Rate	833MHz	1.2GHz	552MHz	450MHz	1GHz	1.5GHz	400MHz	480MHz	900MHz
External Cache	8MB	None	None	8MB	None	None	8MB	8MB	8MB
164.gzip	392	n/a	376	230	545	553	226	165	349
175.vpr	452	n/a	421	285	354	298	384	212	383
176.gcc	617	n/a	577	350	401	588	313	232	500
181.mcf	441	n/a	384	498	276	473	563	356	474
186.crafty	694	n/a	472	304	523	497	334	175	439
197.parser	360	n/a	361	171	362	472	283	211	412
252.eon	645	n/a	395	280	615	650	360	209	465
253.perlbmk	526	n/a	406	215	614	703	246	247	457
254.gap	365	n/a	229	256	443	708	204	171	300
255.vortex	673	n/a	764	312	717	735	294	304	581
256.bzip2	560	n/a	349	258	396	420	334	237	500
300.twolf	658	n/a	479	414	394	403	451	243	473
SPECint_base2000	518	n/a	417	286	454	524	320	225	438
168.wupside	529	360	340	360	416	759	280	284	497
171.swim	1,156	506	761	279	493	1,244	300	285	752
172.mgrid	580	272	462	319	274	558	231	226	377
173.applu	424	298	563	327	280	641	237	150	221
177.mesa	713	302	300	330	541	553	289	273	469
178.galgel	558	468	569	429	335	537	989	735	1,266
179.art	1,540	213	419	969	410	514	995	920	990
183.equake	231	236	347	560	249	739	222	149	211
187.facerec	822	411	258	257	307	451	411	459	718
188.amp	488	221	376	326	294	366	373	313	421
189.lucas	731	237	370	284	349	764	259	205	204
191.fma3d	528	365	302	340	297	427	192	207	302
200.sixtrack	340	256	286	234	170	257	199	159	273
301.aspi	553	278	523	349	371	427	252	189	340
SPECfp_base2000	590	304	400	356	325	549	319	274	427

## **Dynamic Scheduling in Intel Pentium II, III**

**Q: How pipeline 1 to 17 byte 80x86 instructions?**


- ☐ Pentium does not pipeline 80x86 instructions
- ☐ Its decode unit translates the Intel instructions into 72-bit micro-operations (~ MIPS)
- ☐ Sends micro-operations to reorder buffer & reservation stations
- ☐ Many instructions translate to 1 to 4 micro-operations
- ☐ Complex 80x86 instructions are executed by a conventional microprogram (8K x 72 bits) that issues long sequences of micro-operations
- ☐ 14 clocks in total pipeline (~ 3 state machines)

## **Dynamic Scheduling in Intel Pentium**

<u>Parameter</u>	<u>80x86 micro-ops</u>	
Max. instructions issued/clock	3	6
Max. instr. complete exec./clock		5
Max. instr. committed/clock		3
Window (Instrs in reorder buffer)		40
Number of reservations stations	20	
Number of rename registers	40	
No. integer functional units (FUs)	2	
No. floating point FUs	1	
No. SIMD Fl. Pt. FUs	1	
No. memory Fus	1 load + 1 store	

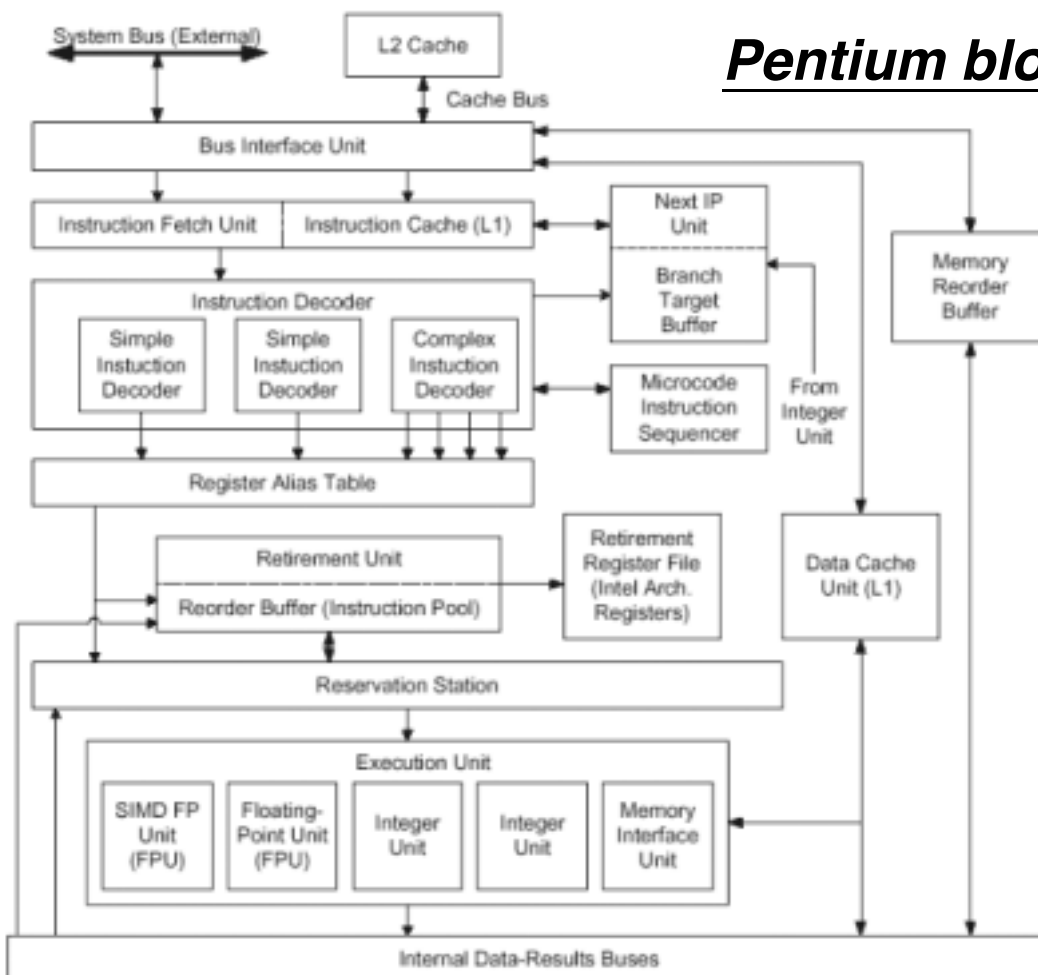
# Pipelining in Pentium

- ❑ 14 clocks in total (~3 state machines)
- ❑ 8 stages are used for in-order instruction fetch, decode, and issue
  - Takes 1 clock cycle to determine length of 80x86 instructions + 2 more to create the micro-operations (uops)
- ❑ 3 stages are used for out-of-order execution in one of 5 separate functional units
- ❑ 3 stages are used for instruction commit



© by Tien-Fu Chen@CCU

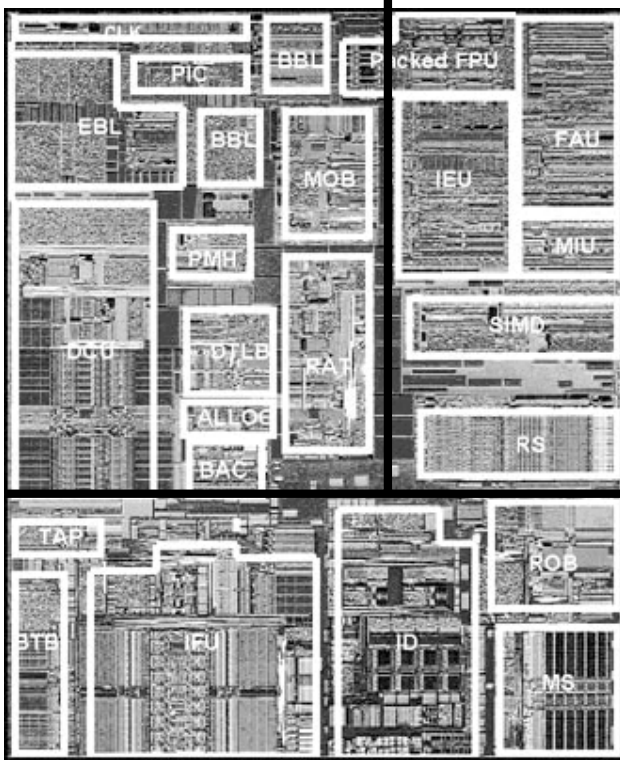
chap4-52



## Pentium block diagram

chap4-53

# **Pentium III Die Photo**



- ☐ EBL/BBL - Bus logic, Front, Back
- ☐ MOB - Memory Order Buffer
- ☐ Picked FPU - MMX Fl. Pt. (SSE)
- ☐ IEU - Integer Execution Unit
- ☐ FAU - Fl. Pt. Arithmetic Unit
- ☐ MIU - Memory Interface Unit
- ☐ DCU - Data Cache Unit
- ☐ PMH - Page Miss Handler
- ☐ DTLB - Data TLB
- ☐ BAC - Branch Address Calculator
- ☐ RAT - Register Alias Table
- ☐ SIMD - Packed Fl. Pt.
- ☐ RS - Reservation Station

---

- ☐ BTB - Branch Target Buffer
- ☐ IFU - Instruction Fetch Unit (+IS)
- ☐ ID - Instruction Decode
- ☐ ROB - Reorder Buffer
- ☐ MS - Micro-instruction Sequencer