

18-447

Computer Architecture

Lecture 1: Introduction and Basics

Prof. Onur Mutlu

Carnegie Mellon University

Spring 2015, 1/12/2015


Major High-Level Goals of This Course

- Understand the principles
- Understand the precedents
- Based on such understanding:
 - Enable you to evaluate tradeoffs of different designs and ideas
 - Enable you to develop principled designs
 - Enable you to develop novel, out-of-the-box designs
- The focus is on:
 - Principles, precedents, and how to use them for new designs
- In Computer Architecture

Role of the (Computer) Architect

Role of the Architect

- Look Backward (Examine old code)***
- Look forward (Listen to the dreamers)***
- Look Up (Nature of the problems)***
- Look Down (Predict the future of technology)***



from Yale Patt's lecture notes

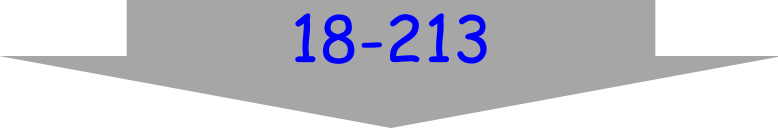
Role of The (Computer) Architect

- Look backward (to the past)
 - Understand tradeoffs and designs, upsides/downsides, past workloads. Analyze and evaluate the past.
- Look forward (to the future)
 - Be the dreamer and create new designs. Listen to dreamers.
 - Push the state of the art. Evaluate new design choices.
- Look up (towards problems in the computing stack)
 - Understand important problems and their nature.
 - Develop architectures and ideas to solve important problems.
- Look down (towards device/circuit technology)
 - Understand the capabilities of the underlying technology.
 - Predict and adapt to the future of technology (you are designing for N years ahead). Enable the future technology.

Takeaways

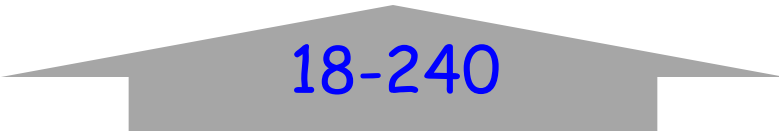
- Being an architect is not easy
- You need to consider **many** things in designing a new system + have good intuition/insight into ideas/tradeoffs
- But, it is fun and can be very technically rewarding
- And, enables a great future
 - E.g., many scientific and everyday-life innovations would not have been possible without architectural innovation that enabled very high performance systems
 - E.g., your mobile phones
- This course will teach you how to become a good computer architect

So, I Hope You Are Here for This



18-213

- How does an assembly program end up executing as digital logic?
- What happens in-between?
- How is a computer designed using logic gates and wires to satisfy specific goals?



18-240

“C” as a model of computation

Programmer’s view of how a computer system works

*Architect/microarchitect’s view:
How to design a computer that meets system design goals.*


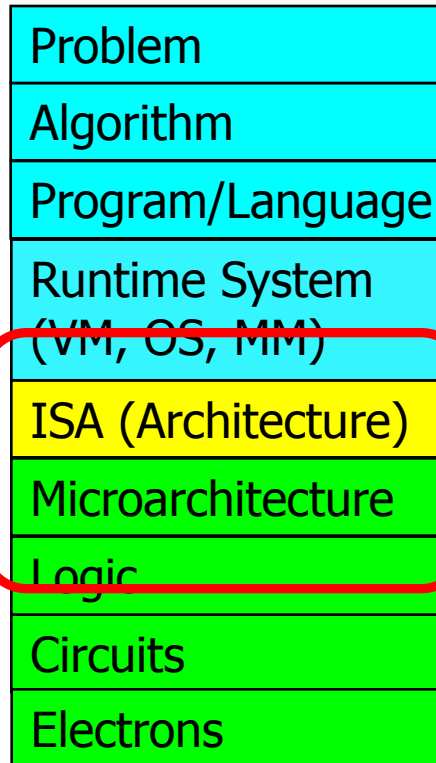
Choices critically affect both the SW programmer and the HW designer

HW designer’s view of how a computer system works

Digital logic as a model of computation

Levels of Transformation

“The purpose of computing is insight” (*Richard Hamming*)
We gain and generate insight by solving problems
How do we ensure problems are solved by electrons?



Aside: A Paper By Hamming

- Hamming, “Error Detecting and Error Correcting Codes,” Bell System Technical Journal 1950.
- Introduced the concept of Hamming distance
 - number of locations in which the corresponding symbols of two equal-length strings is different
- Developed a theory of codes used for error detection and correction
- Also:
- Hamming, “You and Your Research,” Talk at Bell Labs, 1986.
 - <http://www.cs.virginia.edu/~robins/YouAndYourResearch.html>

The Power of Abstraction

■ Levels of transformation create abstractions

- Abstraction: A higher level only needs to know about the interface to the lower level, not how the lower level is implemented
- E.g., high-level language programmer does not really need to know what the ISA is and how a computer executes instructions

■ Abstraction improves productivity

- No need to worry about decisions made in underlying levels
- E.g., programming in Java vs. C vs. assembly vs. binary vs. by specifying control signals of each transistor every cycle

■ Then, why would you want to know what goes on underneath or above?

Crossing the Abstraction Layers

- As long as everything goes well, not knowing what happens in the underlying level (or above) is not a problem.

- What if
 - The program you wrote is running slow?
 - The program you wrote does not run correctly?
 - The program you wrote consumes too much energy?

- What if
 - The hardware you designed is too hard to program?
 - The hardware you designed is too slow because it does not provide the right primitives to the software?

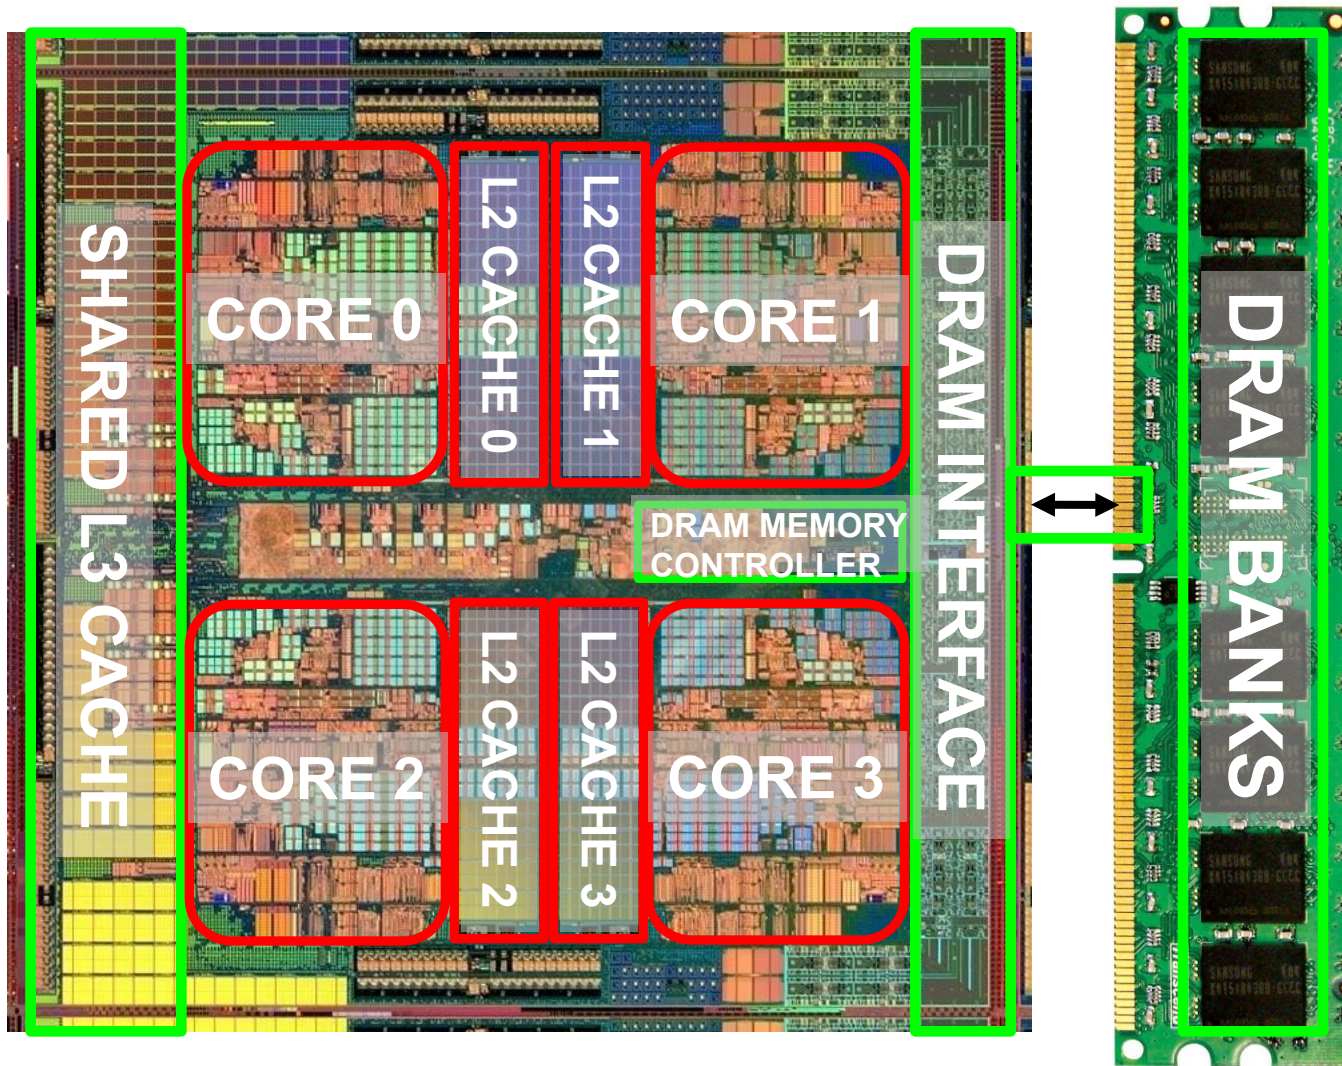
- What if
 - You want to design a much more efficient and higher performance system?

Crossing the Abstraction Layers

- Two key goals of this course are
 - to understand how a processor works underneath the software layer and how decisions made in hardware affect the software/programmer
 - to enable you to be comfortable in making design and optimization decisions that cross the boundaries of different layers and system components


An Example: Multi-Core Systems

Multi-Core
Chip



*Die photo credit: AMD Barcelona

Unexpected Slowdowns in Multi-Core

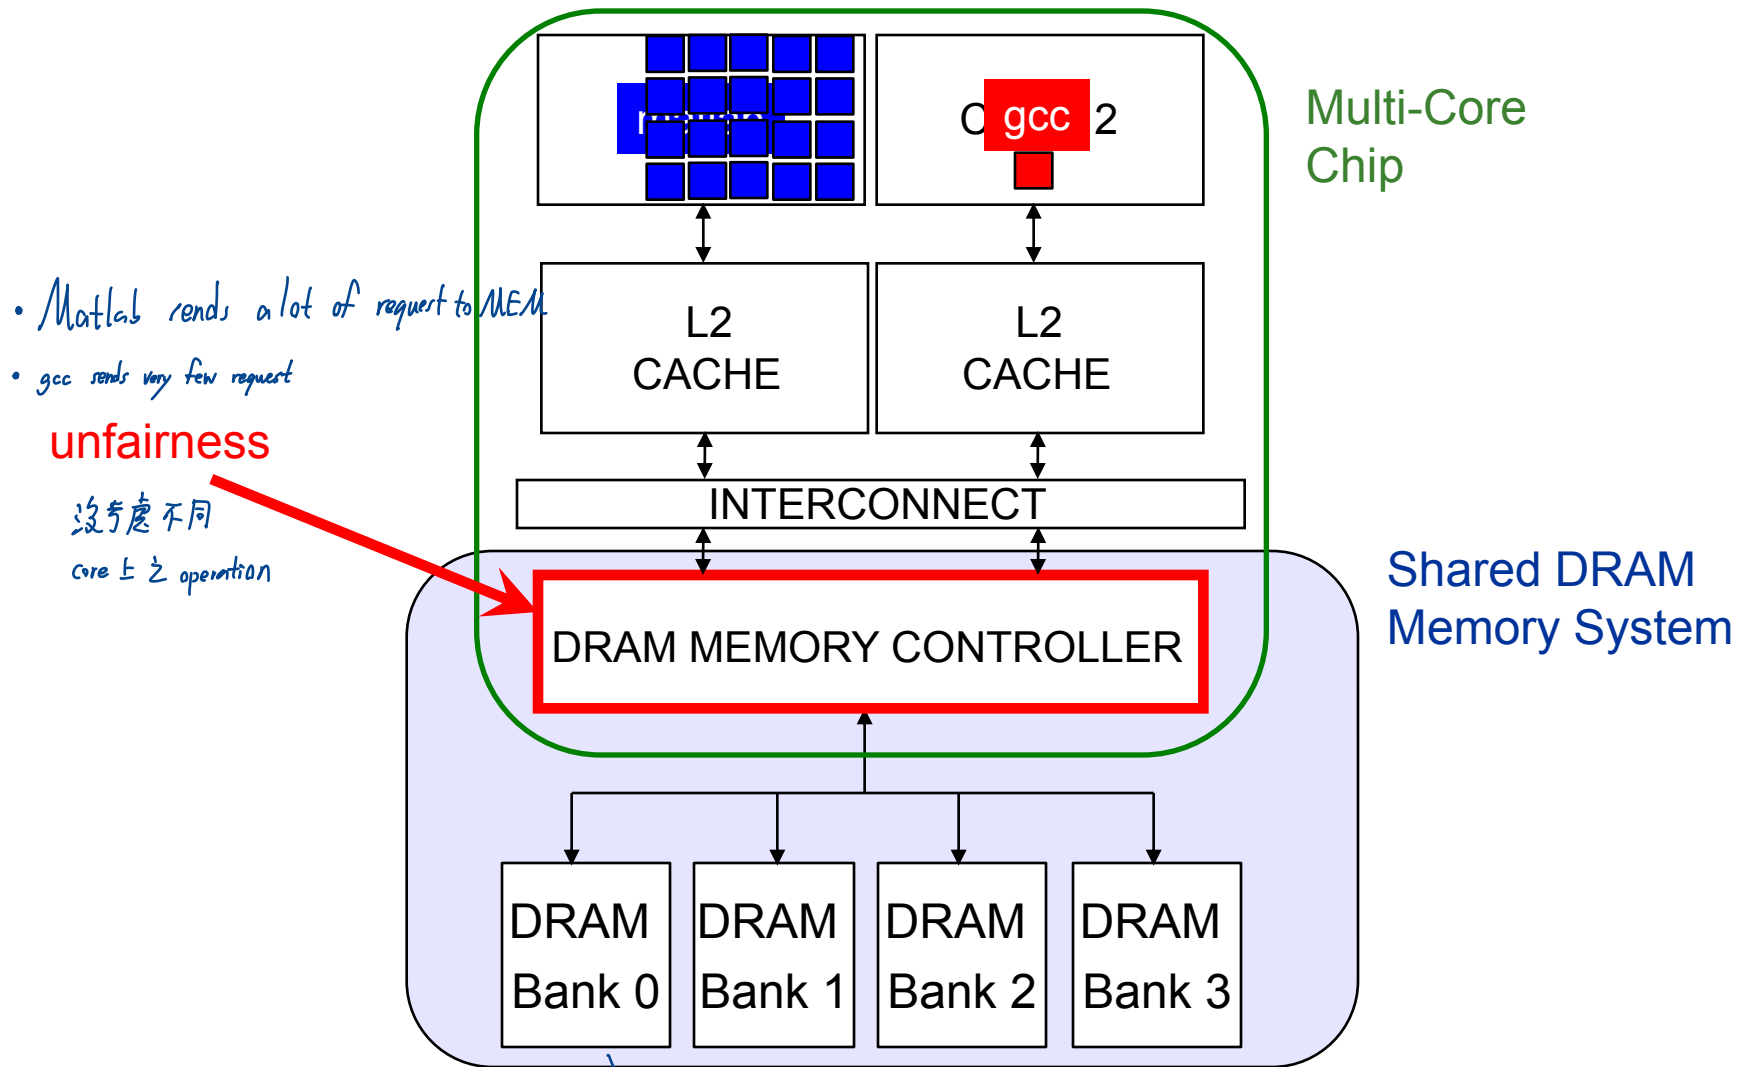


Moscibroda and Mutlu, “[Memory performance attacks: Denial of memory service in multi-core systems](#),” USENIX Security 2007.

A Question or Two

- Can you figure out why there is a disparity in slowdowns if you do not know how the system executes the programs?
 - *Matlab has much better cache locality*
- Can you fix the problem without knowing what is happening “underneath”?

Why the Disparity in Slowdowns?



DRAM Bank Operation

Access Address:

(Row 0, Column 0)

(Row 0, Column 1) *hit the row buffer!*

(Row 0, Column 85) *hit the row buffer!*

(Row 1, Column 0) *conflict!*

Row address 0

Row decoder

Columns

Rows

Row 1

Row Buffer **CONFLICT!**

Column address 05

Column mux

Data

2D DRAM access 中

① 先依 row num 至 row decoder 取出-整个 row 至 row buffer

② 再用 col num 至 col decoder 取出 (row, col) 之 data

DRAM Controllers

- A row-conflict memory access takes significantly longer than a row-hit access
- Current controllers take advantage of the row buffer
- Commonly used scheduling policy (FR-FCFS) [Rixner 2000]*
 - (1) Row-hit first: Service row-hit memory accesses first *⇒ optimize row-hit rate*
 - (2) Oldest-first: Then service older accesses first
- This scheduling policy aims to maximize DRAM throughput

*Rixner et al., “Memory Access Scheduling,” ISCA 2000.

*Zuravleff and Robinson, “Controller for a synchronous DRAM ...,” US Patent 5,630,096, May 1997.

The Problem

- Multiple applications share the DRAM controller
- DRAM controllers designed to maximize DRAM data throughput

造成某些 process starvation

- DRAM scheduling policies are unfair to some applications
 - Row-hit first: unfairly prioritizes apps with high row buffer locality
 - Threads that keep on accessing the same row e.g. Matlab > gcc
 - Oldest-first: unfairly prioritizes memory-intensive applications
requires a large amount of MEM
- DRAM controller vulnerable to denial of service attacks
 - Can write programs to exploit unfairness

A Memory Performance Hog

```
// initialize large arrays A, B

for (j=0; j<N; j++) {
    index = j*linesize; streaming
    A[index] = B[index];
    ...
}
```

STREAM

- Sequential memory access
- Very high row buffer locality (96% hit rate)
- Memory intensive

```
// initialize large arrays A, B


for (j=0; j<N; j++) {
    index = rand(); random
    A[index] = B[index];
    ...
}
```

RANDOM

- Random memory access
- **Very low row buffer locality** (3% hit rate)
- Similarly memory intensive

Moscibroda and Mutlu, “[Memory Performance Attacks](#),” USENIX Security 2007.

What Does the Memory Hog Do?



Row size: 8KB, cache block size: 64B
128 (8KB/64B) requests of T0 serviced before T1

Moscibroda and Mutlu, “[Memory Performance Attacks](#),” USENIX Security 2007.

Now That We Know What Happens Underneath

- How would you solve the problem?

∴ MEM 具 row buffer ⇒ 降低 row-buffer miss ⇒ prioritize row buffer locality 高 2 thread

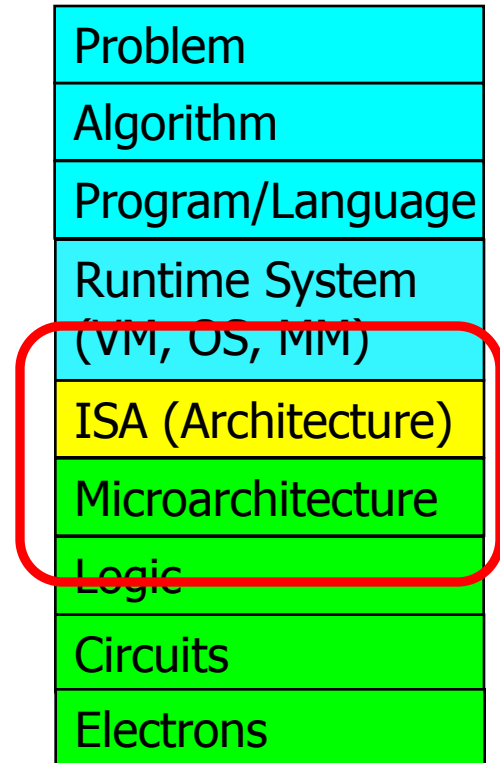
unfair
↑

- What is the right place to solve the problem?

- Programmer?
- System software?
- Compiler? ⇒ has no idea for the other program
- Hardware (Memory controller)?
- Hardware (DRAM)?
- Circuits?

- Two other goals of this course:

- Enable you to think critically
- Enable you to think broadly



Reading on Memory Performance Attacks

- Thomas Moscibroda and Onur Mutlu,
"Memory Performance Attacks: Denial of Memory Service in Multi-Core Systems"
*Proceedings of the 16th USENIX Security Symposium (**USENIX SECURITY**),*
pages 257-274, Boston, MA, August 2007. [Slides \(ppt\)](#)
- One potential reading for your Homework 1 assignment

If You Are Interested ... Further Readings

- Onur Mutlu and Thomas Moscibroda,
"Stall-Time Fair Memory Access Scheduling for Chip Multiprocessors"
*Proceedings of the 40th International Symposium on Microarchitecture (**MICRO**), pages 146-158, Chicago, IL, December 2007. Slides (ppt)*
- Sai Prashanth Muralidhara, Lavanya Subramanian, Onur Mutlu, Mahmut Kandemir, and Thomas Moscibroda,
"Reducing Memory Interference in Multicore Systems via Application-Aware Memory Channel Partitioning"
*Proceedings of the 44th International Symposium on Microarchitecture (**MICRO**), Porto Alegre, Brazil, December 2011. Slides (pptx)*

Takeaway

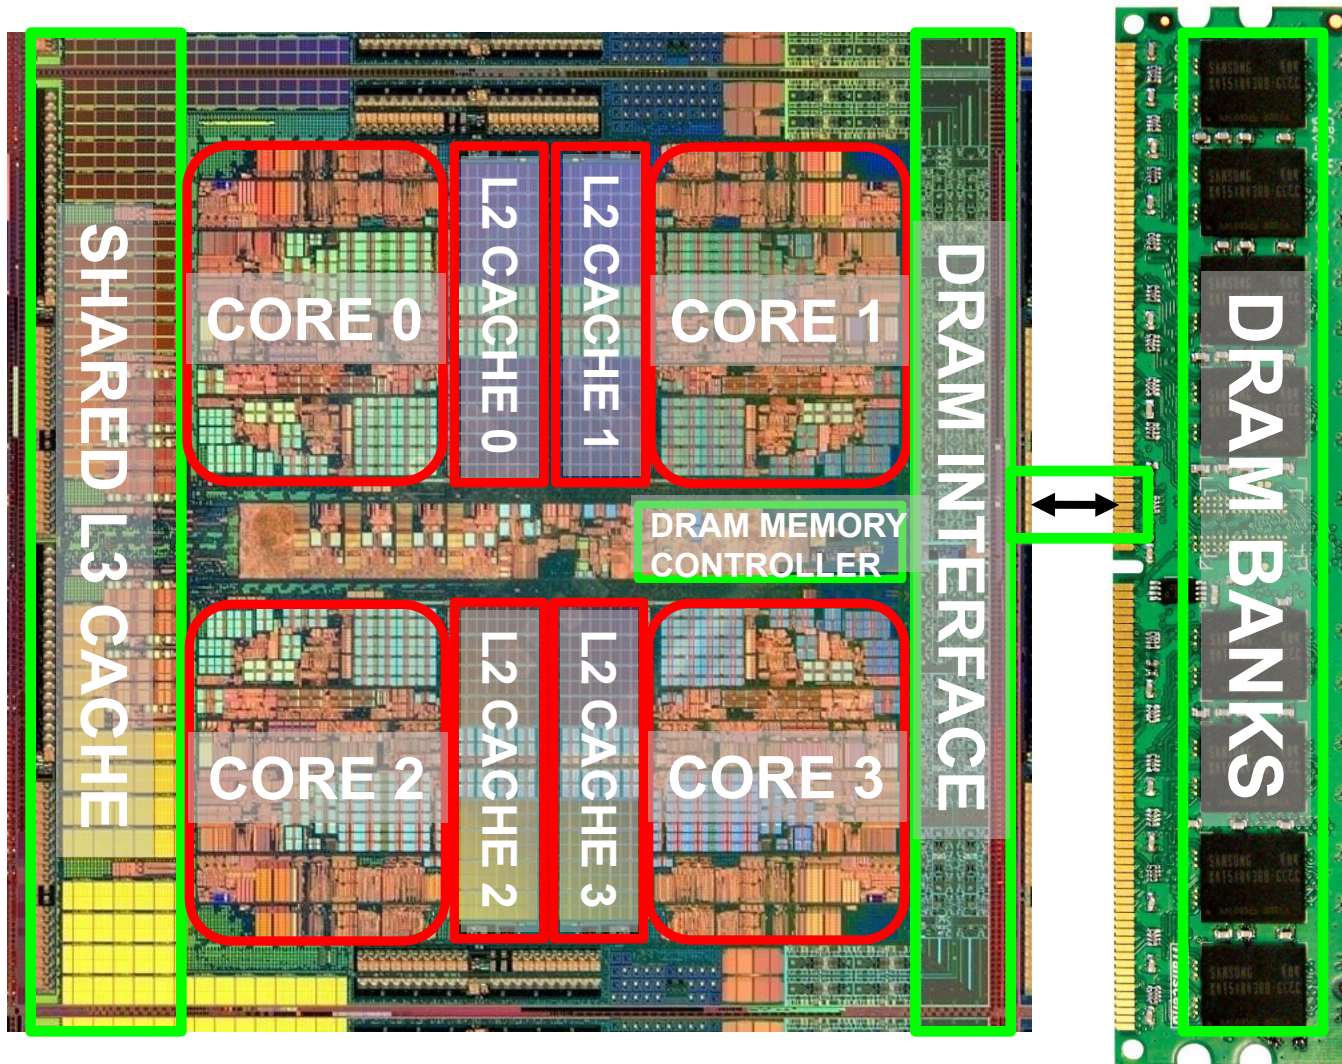
- Breaking the abstraction layers (between components and transformation hierarchy levels) and knowing what is underneath enables you to solve problems

Another Example

- DRAM Refresh

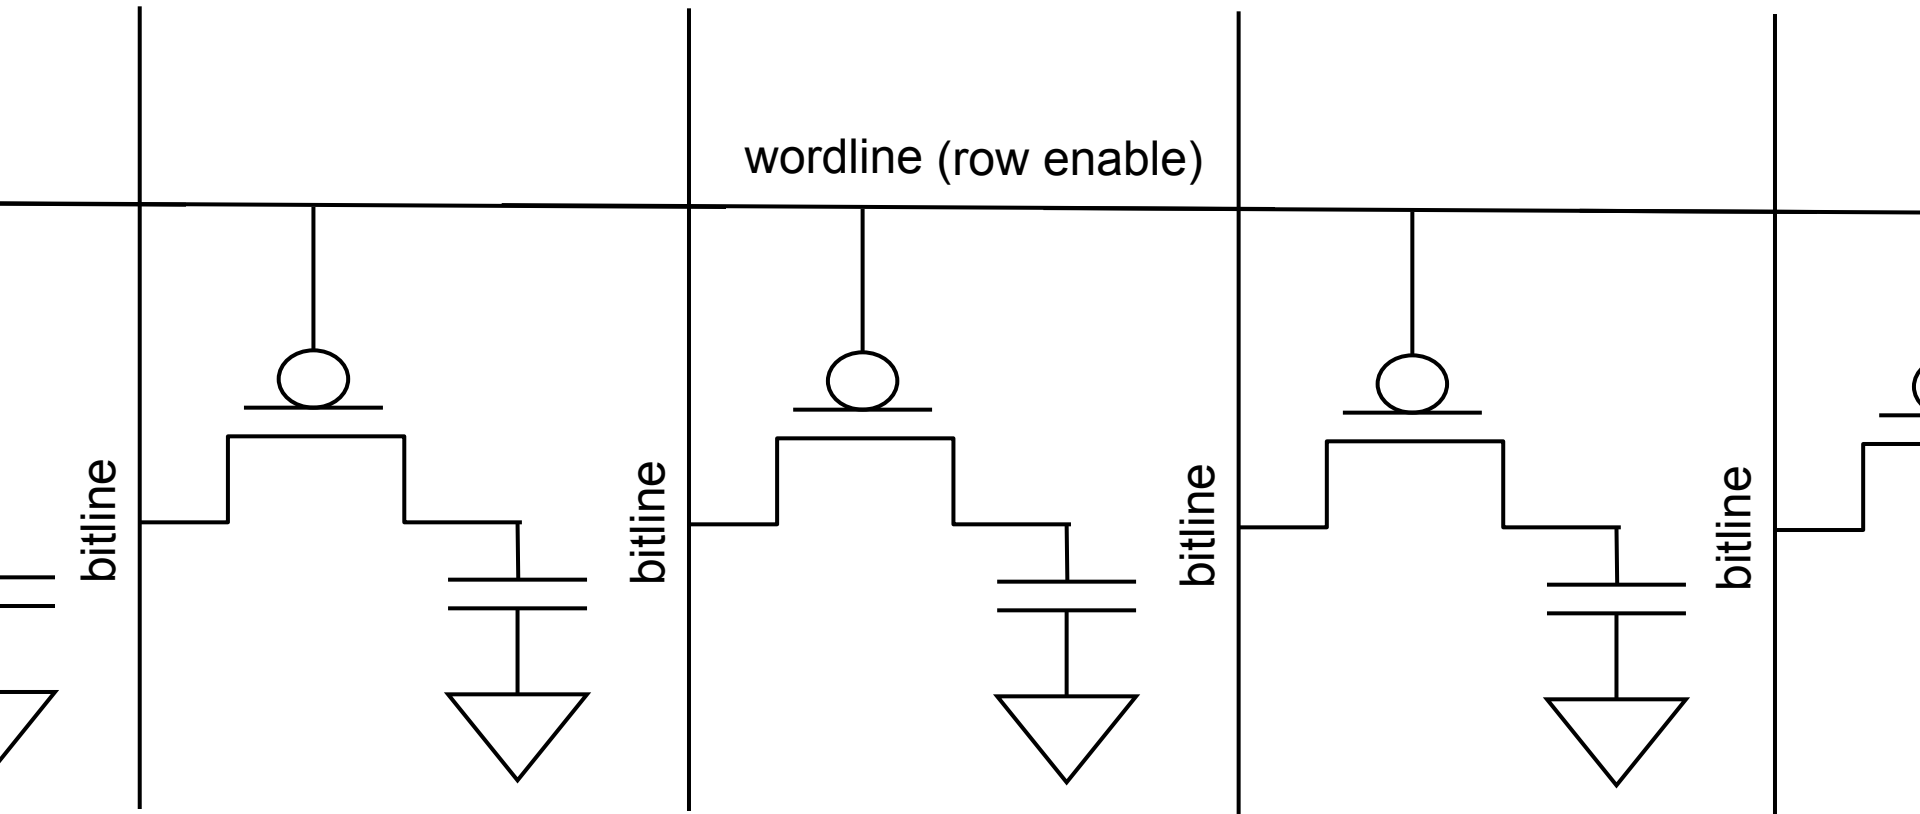
DRAM in the System

Multi-Core
Chip



*Die photo credit: AMD Barcelona

A DRAM Cell



- A DRAM cell consists of a capacitor and an access transistor
- It stores data in terms of charge in the capacitor
- A DRAM chip consists of (10s of 1000s of) rows of such cells


DRAM Refresh

- DRAM capacitor charge leaks over time
- The memory controller needs to refresh each row periodically to restore charge
 - Activate each row every N ms
 - Typical $N = 64$ ms
- Downsides of refresh
 - Energy consumption: Each refresh consumes energy
 - Performance degradation: DRAM rank/bank unavailable while refreshed
 - QoS/predictability impact: (Long) pause times during refresh
 - Refresh rate limits DRAM capacity scaling


First, Some Analysis

- Imagine a system with 1 ExaByte DRAM
- Assume a row size of 8 KiloBytes
- How many rows are there?
- How many refreshes happen in 64ms?
- What is the total power consumption of DRAM refresh?
- What is the total energy consumption of DRAM refresh during a day?
- Part of your Homework 1

Refresh Overhead: Performance



Refresh Overhead: Energy




How Do We Solve the Problem?

- Do we need to refresh all rows every 64ms?

OS 做 memory management, 提供。那些 memory space 沒在使用的。但要如何 expose 該 information 是問題

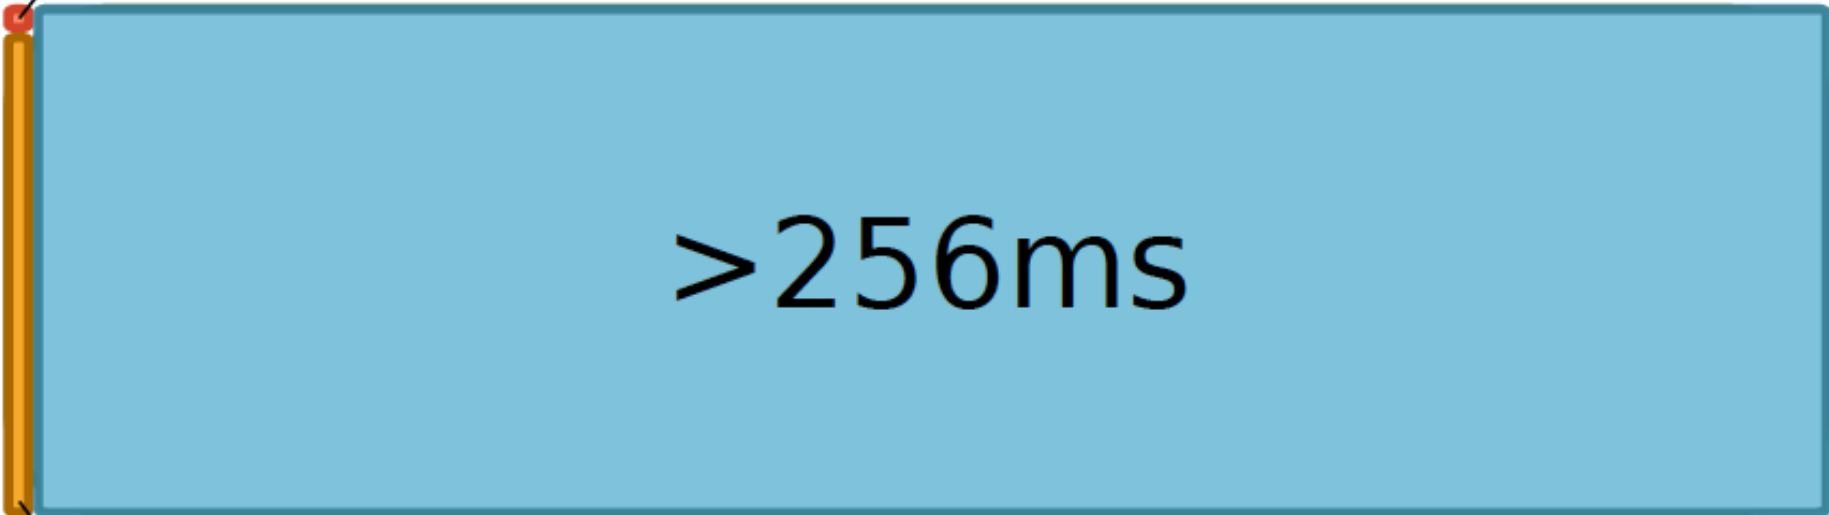
- What if we knew what happened underneath and exposed that information to upper layers?



Underneath: Retention Time Profile of DRAM

64-128ms

不同 DRAM cell 的 Capacity 不同



>256ms

128-256ms

Taking Advantage of This Profile

- Expose this retention time profile information to
 - the memory controller
 - the operating system
 - the programmer?
 - the compiler?
- How much information to expose?
 - Affects hardware/software overhead, power consumption, verification complexity, cost
- How to determine this profile information?
 - Also, who determines it?

An Example: RAIDR

■ Observation: Most DRAM rows can be refreshed much less often without losing data [Kim+, EDL'09][Liu+ ISCA'13]

■ Key idea: Refresh rows containing weak cells more frequently, other rows less frequently

1. **Profiling:** Profile retention time of all rows

2. **Binning:** Store rows into bins by retention time in memory controller

Efficient storage with Bloom Filters (only 1.25KB for 32GB memory)

3. **Refreshing:** Memory controller refreshes rows in different bins at different rates



■ Results: 8-core, 32GB, SPEC, TPC-C, TPC-H

□ 74.6% refresh reduction @ 1.25KB storage

□ ~16%/20% DRAM dynamic/idle power reduction

□ ~9% performance improvement

□ Benefits increase with DRAM capacity



Reading on RAIDR

- Jamie Liu, Ben Jaiyen, Richard Veras, and Onur Mutlu,
"RAIDR: Retention-Aware Intelligent DRAM Refresh"
Proceedings of the 39th International Symposium on Computer Architecture (ISCA), Portland, OR, June 2012. [Slides \(pdf\)](#)
- One potential reading for your Homework 1 assignment
記錄 retention time profile \Rightarrow expose 給 DRAM controller \Rightarrow DRAM controller 依 該 information 去 refresh

If You Are Interested ... Further Readings

- Onur Mutlu,
"Memory Scaling: A Systems Architecture Perspective"
Technical talk at MemCon 2013 (MEMCON), Santa Clara, CA, August 2013.
[Slides \(pptx\)](#) [\(pdf\)](#) [Video](#)
- Kevin Chang, Donghyuk Lee, Zeshan Chishti, Alaa Alameldeen, Chris Wilkerson, Yoongu Kim, and Onur Mutlu,
"Improving DRAM Performance by Parallelizing Refreshes with Accesses"
Proceedings of the 20th International Symposium on High-Performance Computer Architecture (HPCA), Orlando, FL, February 2014. [Slides \(pptx\)](#) [\(pdf\)](#)

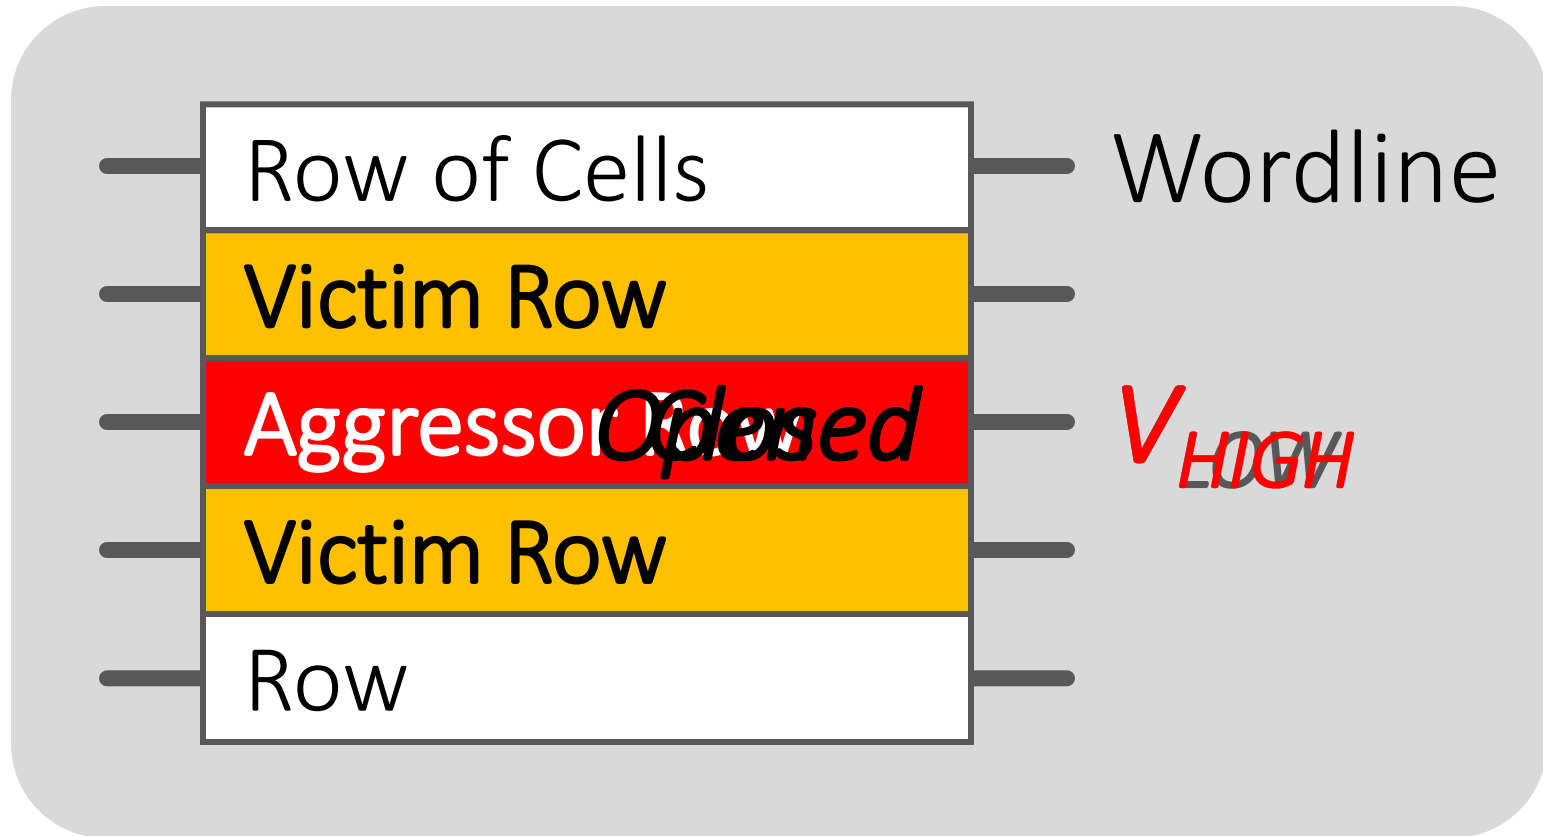
Takeaway

- Breaking the abstraction layers (between components and transformation hierarchy levels) and knowing what is underneath enables you to solve problems and design better future systems
- Cooperation between multiple components and layers can enable more effective solutions and systems

Yet Another Example

- DRAM Row Hammer (or, DRAM Disturbance Errors)


Disturbance Errors in Modern DRAM



Repeatedly opening and closing a row enough times within a refresh interval induces **disturbance errors** in adjacent rows in **most real DRAM chips you can buy today** ~~☆☆☆~~.


Most DRAM Modules Are At Risk

A company




Up to
 1.0×10^7
errors

B company



Up to
 2.7×10^6
errors



C company




Up to
 3.3×10^5
errors

Kim+, "Flipping Bits in Memory Without Accessing Them: An Experimental Study of DRAM Disturbance Errors," ISCA 2014.


x86 CPU





DRAM Module




```
loop:  
  mov  (X), %eax  
  mov  (Y), %ebx  
  clflush (X)  
  clflush (Y)  
  mfence  
  jmp  loop
```




x86 CPU





DRAM Module




```
loop:  
  mov  (X), %eax  
  mov  (Y), %ebx  
  clflush (X)  
  clflush (Y)  
  mfence  
  jmp  loop
```




x86 CPU





DRAM Module




```
loop:  
  mov  (X), %eax  
  mov  (Y), %ebx  
  clflush (X)  
  clflush (Y)  
  mfence  
  jmp  loop
```



x86 CPU




DRAM Module



loop:

```
mov  (X), %eax  
mov  (Y), %ebx  
clflush (X)  
clflush (Y)  
mfence  
jmp  loop
```




Observed Errors in Real Systems

CPU Architecture	Errors	Access-Rate
Intel Haswell (2013)	22.9K	12.3M/sec
Intel Ivy Bridge (2012)	20.7K	11.7M/sec
Intel Sandy Bridge (2011)	16.1K	11.6M/sec
AMD Piledriver (2012)	59	6.1M/sec

- *A real reliability & security issue*
- *In a more controlled environment, we can induce as many as **ten million** disturbance errors*

Errors *vs.* Vintage



All modules from 2012-2013 are vulnerable

How Do We Solve The Problem?

- **Do business as usual but better:** Improve circuit and device technology such that disturbance does not happen.

Use stronger error correcting codes. *Using redundancy*

- **Tolerate it:** Make DRAM and controllers more intelligent so that they can proactively fix the errors
- **Eliminate or minimize it:** Replace DRAM with a different technology that does not have the problem
- **Embrace it:** Design heterogeneous-reliability memories that map error-tolerant data to less reliable portions

- ...

More on DRAM Disturbance Errors

- Yoongu Kim, Ross Daly, Jeremie Kim, Chris Fallin, Ji Hye Lee, Donghyuk Lee, Chris Wilkerson, Konrad Lai, and Onur Mutlu,
"Flipping Bits in Memory Without Accessing Them: An Experimental Study of DRAM Disturbance Errors"
Proceedings of the 41st International Symposium on Computer Architecture (ISCA), Minneapolis, MN, June 2014. [Slides \(pptx\)](#) [\(pdf\)](#)
[Lightning Session Slides \(pptx\)](#) [\(pdf\)](#) [Source Code and Data](#)
- Source Code to Induce Errors in Modern DRAM Chips
 - <https://github.com/CMU-SAFARI/rowhammer>
- One potential reading for your Homework 1 assignment

Recap: Some Goals of 447

- Teach/enable/empower you to:
 - ❑ Understand how a computing platform (processor + memory + interconnect) works
 - ❑ Implement a simple platform (with not so simple parts), with a focus on the processor and memory
 - ❑ Understand how decisions made in hardware affect the software/programmer as well as hardware designer
 - ❑ Think critically (in solving problems)
 - ❑ Think broadly across the levels of transformation
 - ❑ Understand how to analyze and make tradeoffs in design

Review: Major High-Level Goals of This Course

- Understand the principles
- Understand the precedents
- Based on such understanding:
 - Enable you to evaluate tradeoffs of different designs and ideas
 - Enable you to develop principled designs
 - Enable you to develop novel, out-of-the-box designs
- The focus is on:
 - Principles, precedents, and how to use them for new designs
- In Computer Architecture

A Note on Hardware vs. Software

- This course is classified under “Computer Hardware”
- However, you will be much more capable if you master both hardware and software (and the interface between them)
 - Can develop better software if you understand the underlying hardware
 - Can design better hardware if you understand what software it will execute
 - Can design a better computing system if you understand both
- This course covers the HW/SW interface and microarchitecture
 - We will focus on tradeoffs and how they affect software

Required Readings for This Week

- Patt, “Requirements, Bottlenecks, and Good Fortune: Agents for Microprocessor Evolution,” Proceedings of the IEEE 2001.
- One of
 - Moscibroda and Mutlu, “Memory Performance Attacks: Denial of Memory Service in Multi-Core Systems,” USENIX Security 2007.
 - Liu+, “RAIDR: Retention-Aware Intelligent DRAM Refresh,” ISCA 2012.
 - Kim+, “Flipping Bits in Memory Without Accessing Them: An Experimental Study of DRAM Disturbance Errors,” ISCA 2014.
- P&P Chapter 1 (Fundamentals)
- P&H Chapters 1 and 2 (Intro, Abstractions, ISA, MIPS)
- Reference material throughout the course
 - MIPS ISA Reference Manual + x86 ISA Reference Manual
 - <http://www.ece.cmu.edu/~ece447/s15/doku.php?id=techdocs>

A Note on Books

- None required
- But, I expect you to be resourceful in finding and doing the readings...

Recitations Next Week

- MIPS ISA Tutorial
 - You can attend any recitation session

18-447

Computer Architecture

Lecture 1: Introduction and Basics

Prof. Onur Mutlu

Carnegie Mellon University

Spring 2015, 1/12/2015

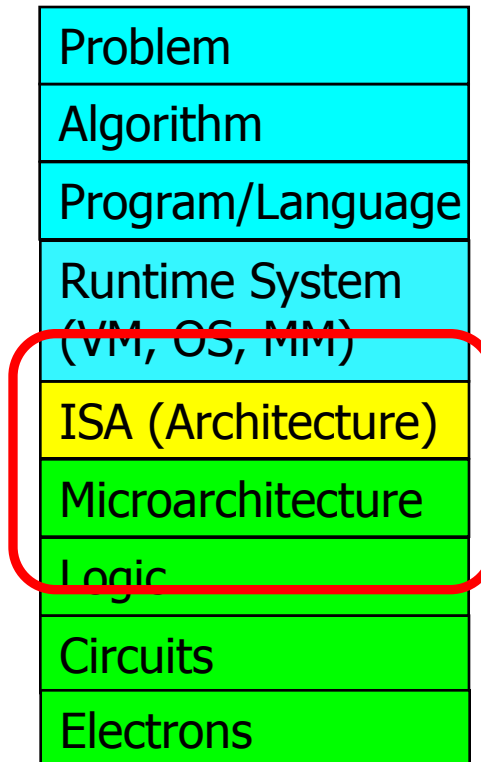
What Will You Learn

- **Computer Architecture:** The science and art of designing, selecting, and interconnecting hardware components and designing the hardware/software interface to create a computing system that meets functional, performance, energy consumption, cost, and other specific goals.
- **Traditional definition:** “The term *architecture* is used here to describe the attributes of a computer system as seen by the programmer, i.e., the conceptual structure and behavior as distinct from the organization of the hardware and controls, the logic design, and the physical implementation.” *Gene Amdahl, IBM Systems Journal*, 9:3, 1964



Dr. Amdahl holding a 100gate LSI air-cooled chip. On his desk is a circuit board with the chips on it. This circuit board was for an Amdahl 470 V/6 (photograph dated March 1973).

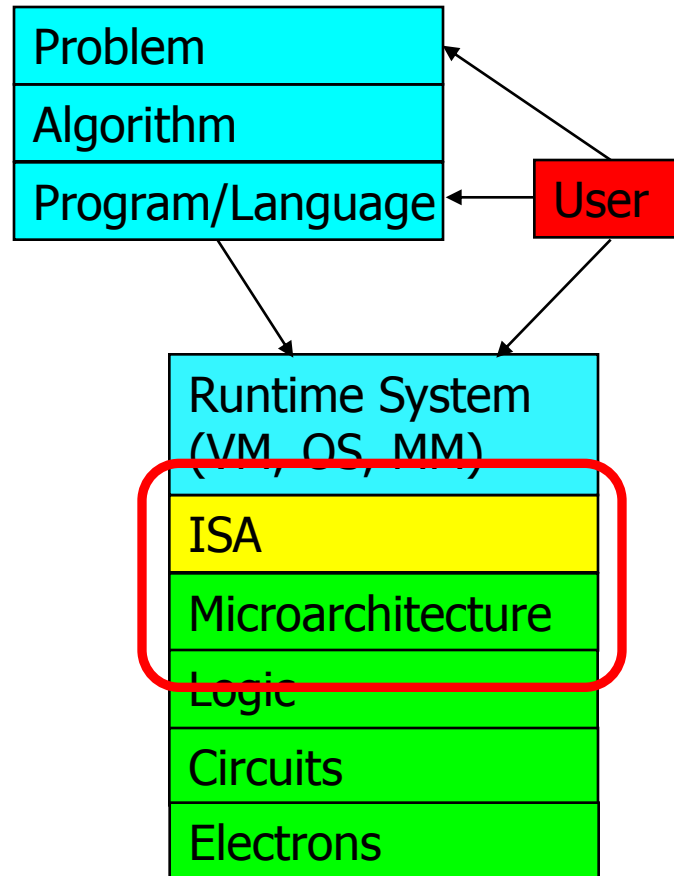
Computer Architecture in Levels of Transformation



- Read: Patt, "Requirements, Bottlenecks, and Good Fortune: Agents for Microprocessor Evolution," Proceedings of the IEEE 2001.

Levels of Transformation, Revisited

- A user-centric view: computer designed for users



- The entire stack should be optimized for user

What Will You Learn?

- Fundamental principles and tradeoffs in designing the hardware/software interface and major components of a modern programmable microprocessor
 - Focus on state-of-the-art (and some recent research and trends)
 - Trade-offs and how to make them
- How to design, implement, and evaluate a functional modern processor
 - Semester-long lab assignments
 - A combination of RTL implementation and higher-level simulation
 - Focus is functionality first (some on “how to do even better”)
- How to dig out information, think critically and broadly
- How to work even harder!

Course Goals

- Goal 1: To familiarize those interested in computer system design with both fundamental operation principles and design tradeoffs of processor, memory, and platform architectures in today's systems.
 - Strong emphasis on fundamentals and design tradeoffs.
- Goal 2: To provide the necessary background and experience to design, implement, and evaluate a modern processor by performing hands-on RTL and C-level implementation.
 - Strong emphasis on functionality and hands-on design.